

ML_SP22_Project_1 (DIY Decision Tree)

Due Date: 4/15 23:59 pm

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```


```
In [2]: df = pd.read_csv('breast_cancer.csv')
X = df.drop(['diagnosis'], axis=1).to_numpy()
# B is benign and is encoded as 1, M is malignant and is encoded as 0
y = df['diagnosis'].apply(lambda x: 0 if x == 'M' else 1).to_numpy()
```

```
In [3]: df.describe()
```

```
Out[3]:
```


	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.048919	0.053470
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	0.038803	0.046371
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	0.000000	0.019161
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	0.020310	0.037461
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	0.033500	0.047581
75%	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	0.074000	0.065119
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	0.201200	0.173401

8 rows × 10 columns



```
In [4]: sns.countplot(x='diagnosis',data=df, palette='RdBu_r')
```

```
Out[4]: <AxesSubplot:xlabel='diagnosis', ylabel='count'>
```



```
In [5]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

First build the model with the standard sklearn library

```
In [6]: from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier(max_depth=10)
model.fit(X_train, y_train)
```

```
Out[6]: DecisionTreeClassifier(max_depth=10)
```

```
In [7]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
predictions = model.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
print(accuracy_score(y_test, predictions))
```

```
[[38  4]
 [ 2 70]]
```

	precision	recall	f1-score	support
0	0.95	0.90	0.93	42
1	0.95	0.97	0.96	72
accuracy			0.95	114
macro avg	0.95	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114

0.9473684210526315

Second use the implementation of the blog to build the model

<https://towardsdatascience.com/implementing-a-decision-tree-from-scratch-f5358ff9c4bb>

In [8]:

```
from DT_orig import DecisionTree
model = DecisionTree(max_depth=10)
model.fit(X_train, y_train)
```

Done fitting

In [9]:

```
from DT_orig import accuracy_score
predictions = model.predict(X_test)
print(accuracy_score(y_test, predictions))
```

0.956140350877193

Note that the original implementation will not work if y is a categorical variable and it is expecting numpy array instead of DataFrame

In [10]:

```
X = df.drop(['diagnosis'], axis=1).to_numpy()
#y = df['diagnosis'].apply(lambda x: 0 if x == 'M' else 1).to_numpy()
y = df['diagnosis'].to_numpy()
y[:10]
```

Out[10]:

array(['M', 'M', 'M', 'M', 'M', 'M', 'M', 'M', 'M', 'M'], dtype=object)

In [19]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
model.fit(X_train, y_train)
print(accuracy_score(y_test, predictions))
```

```

-----
KeyError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11464\1472857777.py in <module>
      1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
----> 2 model.fit(X_train, y_train)
      3 print(accuracy_score(y_test, predictions))

~\Downloads\python\DT_Iman_Toussi.py in fit(self, X, y)
     35         # call the _fit method
     36         x = X.to_numpy()
----> 37         self._fit(x, y)
     38         # end TODO
     39         print("Done fitting")

~\Downloads\python\DT_Iman_Toussi.py in _fit(self, X, y)
     48
     49     def _fit(self, X, y):
----> 50         self.root = self._build_tree(X, y)
     51
     52     def _predict(self, X):

~\Downloads\python\DT_Iman_Toussi.py in _build_tree(self, X, y, depth)
     81         # get best split
     82         rnd_feats = np.random.choice(self.n_features, self.n_features, replace=False)
----> 83         best_feat, best_thresh = self._best_split(X, y, rnd_feats)
     84
     85         # grow children recursively

~\Downloads\python\DT_Iman_Toussi.py in _best_split(self, X, y, features)
    144         thresholds = np.unique(X_feat)
    145         for thresh in thresholds:
--> 146             score = self._information_gain(X_feat, y, thresh)
    147
    148             if score > split['score']:

~\Downloads\python\DT_Iman_Toussi.py in _information_gain(self, X, y, thresh)
    128         return 0
    129
--> 130         child_loss = (n_left / n) * self._entropy(y[left_idx]) + (n_right / n) * self._entropy(y[right_idx])
    131         child_loss2 = (n_left / n) * self._gini(y[left_idx]) + (n_right / n) * self._gini(y[right_idx])
    132         # end TODO

F:\CODE\anaconda3\lib\site-packages\pandas\core\series.py in __getitem__(self, key)
    964         return self._get_values(key)
    965

```

```

--> 966         return self._get_with(key)
      967
      968     def _get_with(self, key):

F:\CODE\anaconda3\lib\site-packages\pandas\core\series.py in _get_with(self, key)
      999         # (i.e. self.iloc) or label-based (i.e. self.loc)
     1000         if not self.index._should_fallback_to_positional():
-> 1001             return self.loc[key]
     1002         else:
     1003             return self.iloc[key]

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in __getitem__(self, key)
      929
      930         maybe_callable = com.apply_if_callable(key, self.obj)
--> 931         return self._getitem_axis(maybe_callable, axis=axis)
      932
      933     def _is_scalar_access(self, key: tuple):

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _getitem_axis(self, key, axis)
     1151         raise ValueError("Cannot index with multidimensional key")
     1152
-> 1153         return self._getitem_iterable(key, axis=axis)
     1154
     1155         # nested tuple slicing

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _getitem_iterable(self, key, axis)
     1091
     1092         # A collection of keys
-> 1093         keyarr, indexer = self._get_listlike_indexer(key, axis)
     1094         return self.obj._reindex_with_indexers(
     1095             {axis: [keyarr, indexer]}, copy=True, allow_dups=True

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _get_listlike_indexer(self, key, axis)
     1312         keyarr, indexer, new_indexer = ax._reindex_non_unique(keyarr)
     1313
-> 1314         self._validate_read_indexer(keyarr, indexer, axis)
     1315
     1316         if needs_i8_conversion(ax.dtype) or isinstance(

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _validate_read_indexer(self, key, indexer, axis)
     1375
     1376         not_found = list(ensure_index(key)[missing_mask.nonzero()[0]].unique())
-> 1377         raise KeyError(f"{not_found} not in index")
     1378
     1379

```

KeyError: '[62, 65, 111, 207] not in index'

Finally use your own improved implementation to build the model

```
In [12]: from DT_Iman_Toussi import DecisionTreeModel
# replace the above with your version
model = DecisionTreeModel(max_depth=10)


X = df.drop(['diagnosis'], axis=1)
y = df['diagnosis'].apply(lambda x: 0 if x == 'B' else 1)
# make sure your model will work with y being a categorcal variable as well
#y = df['diagnosis']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

```
In [13]: X_train.head()
```

```
Out[13]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_n
408	17.99	20.66	117.80	991.7	0.10360	0.13040	0.120100	0.088240	0.
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.198000	0.104300	0.
307	9.00	14.40	56.36	246.3	0.07005	0.03116	0.003681	0.003472	0.
386	12.21	14.09	78.78	462.0	0.08108	0.07823	0.068390	0.025340	0.
404	12.34	14.95	78.29	469.1	0.08682	0.04571	0.021090	0.020540	0.

5 rows × 30 columns



```
In [14]: y_train.head()
```

```
Out[14]: 408    1
4        1
307     0
386     0
404     0
Name: diagnosis, dtype: int64
```

```
In [15]: type(X_train)
```

```
Out[15]: pandas.core.frame.DataFrame
```

```
In [16]: model.fit(X_train, y_train)
```

```
-----  
KeyError                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_11464\180087699.py in <module>  
----> 1 model.fit(X_train, y_train)  
  
~\Downloads\python\DT_Iman_Toussi.py in fit(self, X, y)  
    35         # call the _fit method  
    36         x = X.to_numpy()  
----> 37         self._fit(x, y)  
    38         # end TODO  
    39         print("Done fitting")  
  
~\Downloads\python\DT_Iman_Toussi.py in _fit(self, X, y)  
    48  
    49     def _fit(self, X, y):  
----> 50         self.root = self._build_tree(X, y)  
    51  
    52     def _predict(self, X):  
  
~\Downloads\python\DT_Iman_Toussi.py in _build_tree(self, X, y, depth)  
    81         # get best split  
    82         rnd_feats = np.random.choice(self.n_features, self.n_features, replace=False)  
----> 83         best_feat, best_thresh = self._best_split(X, y, rnd_feats)  
    84  
    85         # grow children recursively  
  
~\Downloads\python\DT_Iman_Toussi.py in _best_split(self, X, y, features)  
   144         thresholds = np.unique(X_feat)  
   145         for thresh in thresholds:  
--> 146             score = self._information_gain(X_feat, y, thresh)  
   147  
   148             if score > split['score']:  
  
~\Downloads\python\DT_Iman_Toussi.py in _information_gain(self, X, y, thresh)  
   128         return 0  
   129
```

```

--> 130         child_loss = (n_left / n) * self._entropy(y[left_idx]) + (n_right / n) * self._entropy(y[right_idx])
131         child_loss2 = (n_left / n) * self._gini(y[left_idx]) + (n_right / n) * self._gini(y[right_idx])
132         # end TODO

F:\CODE\anaconda3\lib\site-packages\pandas\core\series.py in __getitem__(self, key)
964         return self._get_values(key)
965
--> 966         return self._get_with(key)
967
968     def _get_with(self, key):

F:\CODE\anaconda3\lib\site-packages\pandas\core\series.py in _get_with(self, key)
999         # (i.e. self.iloc) or label-based (i.e. self.loc)
1000         if not self.index._should_fallback_to_positional():
--> 1001             return self.loc[key]
1002         else:
1003             return self.iloc[key]

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in __getitem__(self, key)
929
930         maybe_callable = com.apply_if_callable(key, self.obj)
--> 931         return self._getitem_axis(maybe_callable, axis=axis)
932
933     def _is_scalar_access(self, key: tuple):

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _getitem_axis(self, key, axis)
1151         raise ValueError("Cannot index with multidimensional key")
1152
--> 1153         return self._getitem_iterable(key, axis=axis)
1154
1155         # nested tuple slicing

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _getitem_iterable(self, key, axis)
1091
1092         # A collection of keys
--> 1093         keyarr, indexer = self._get_listlike_indexer(key, axis)
1094         return self.obj._reindex_with_indexers(
1095             {axis: [keyarr, indexer]}, copy=True, allow_dups=True

F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _get_listlike_indexer(self, key, axis)
1312         keyarr, indexer, new_indexer = ax._reindex_non_unique(keyarr)
1313
--> 1314         self._validate_read_indexer(keyarr, indexer, axis)
1315
1316         if needs_i8_conversion(ax.dtype) or isinstance(

```



```
F:\CODE\anaconda3\lib\site-packages\pandas\core\indexing.py in _validate_read_indexer(self, key, indexer, axis)
1372         if use_interval_msg:
1373             key = list(key)
-> 1374         raise KeyError(f"None of [{key}] are in the [{axis_name}]")
1375
1376         not_found = list(ensure_index(key)[missing_mask.nonzero()[0]].unique())
```

KeyError: "None of [Int64Index([65], dtype='int64')] are in the [index]"

Call your own performance report

In [17]:

```
from DT_Iman_Toussi import classification_report, confusion_matrix, accuracy_score
predictions = model.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
print(accuracy_score(y_test, predictions))
```

```
-----
AttributeError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11464\3186627248.py in <module>
      1 from DT_Iman_Toussi import classification_report, confusion_matrix, accuracy_score
----> 2 predictions = model.predict(X_test)
      3 print(confusion_matrix(y_test, predictions))
      4 print(classification_report(y_test, predictions))
      5 print(accuracy_score(y_test, predictions))

~\Downloads\python\DT_Iman_Toussi.py in predict(self, X)
     43     # call the predict method
     44     x = X.to_numpy()
----> 45     self._predict(x)
     46     # return ...
     47     # end TODO

~\Downloads\python\DT_Iman_Toussi.py in _predict(self, X)
     51
     52     def _predict(self, X):
----> 53         predictions = [self._traverse_tree(x, self.root) for x in X]
     54         return np.array(predictions)
     55

~\Downloads\python\DT_Iman_Toussi.py in <listcomp>(.0)
     51
     52     def _predict(self, X):
```

```

---> 53     predictions = [self._traverse_tree(x, self.root) for x in X]
      54     return np.array(predictions)
      55

~\Downloads\python\DT_Iman_Toussi.py in _traverse_tree(self, x, node)
      159     node.
      160     '''
--> 161     if node.is_leaf():
      162         return node.value
      163

```

AttributeError: 'NoneType' object has no attribute 'is_leaf'

Finally call your RandomForest Model just like the standard sklearn library

In [18]:

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

rfc = RandomForestClassifier(n_estimators=100)
rfc.fit(X_train, y_train)
rfc_pred = rfc.predict(X_test)
print(classification_report(y_test, rfc_pred))
print(accuracy_score(y_test, rfc_pred))

```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	72
1	1.00	0.88	0.94	42
accuracy			0.96	114
macro avg	0.97	0.94	0.95	114
weighted avg	0.96	0.96	0.96	114

0.956140350877193

In [44]:

```

# Type your code here
from DT_Iman_Toussi import RandomForestModel
from DT_Iman_Toussi import classification_report, confusion_matrix, accuracy_score

rfc = RandomForestModel(n_estimators=100)
rfc.fit(X_train, y_train)
rfc_pred = rfc.predict(X_test)

```

```
print(classification_report(y_test, rfc_pred))
print(accuracy_score(y_test, rfc_pred))
```

Traceback (most recent call last):

```
File "F:\CODE\anaconda3\lib\site-packages\IPython\core\interactiveshell.py", line 3444, in run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
```

```
File "C:\Users\psy\AppData\Local\Temp\ipykernel_5272\3218034784.py", line 2, in <module>
    from DT_Iman_Toussi import RandomForestModel
```

```
File "C:\Users\psy\Downloads\python\DT_Iman_Toussi.py", line 46
    def _fit(self, X, y):
      ^
```

IndentationError: expected an indented block

For graduate students only, try different value for the impurity threshold for the Decision Tree Model comment on of the impact of the parameter (if there is any) on the model performance

In [22]: *# Type your code here*

In []: